

CS 5740: Project 1

Akash Peri (ap659), Nithish Raja Chidambaram (nr293), Vinayaka Suryanarayana Bommathanahalli (vb247)

Part 1) Unsmoothed ngrams:

Our goal is to compute unsmoothed unigram and bigram probabilities from our corpuses, one negative and one positive. The strategy is to move along the entire corpus with a window of 2, add the current element count as the unigram-count, and then the count of the window as the bigram-count. We can then compute the probability for both on demand.

Preprocessing decisions:

We chose to use the EOS character '\n' to separate documents from each other. After separating documents in this way, we add an explicit start of string character and end of string character - <S> and </S>. We chose not to use '-' as a delimiter; we wanted to treat hyphenated words as a single word. We did, however, use '"' as a delimiter; there were too many edge cases in the corpus to deal with adequately. We did plan on having "'" when used as a conjunction retained as part of the token, but this did not seem to make much of a difference. The words we did choose as delimiters were ',', ';', ':', '!', and '.'. In addition, we experimented with using ',' and '.' as tokens, as they were important parts of speech we did not want to drop. But this created less intelligible sentences and also created additional problems such as unwanted characters like ". ." getting included from the corpus. Based on this, we decided to treat '.' and ',' as delimiters as well.

Part 2) Random sentence generation

When we looked at our sentence generator, we realized that there was a design choice we needed to make, which was to limit sentences to n words. Because we drew the end-of-sentence symbol randomly, we were finding that our sentences were very long. To combat this, we chose to enforce a maximum of n words, with fewer possible if the end-of-sentence symbol is drawn. We used numpy to draw words from our probability distribution. For bigram generation, we maintain a data structure of next possible words that can follow a given word, along with their bigram probabilities.

Unigram sentence generated from pos.txt:

- *of culture advocacy no its are to in is an*
- *both all sense so theme to a illustrated theory by*
- *story that images irony viewer dark best inventive snaps of*

Bigram sentence generated from pos.txt:

- <s> *neither as i love may also a good fight scenes*
- <s> *this is the right questions of the multi-layers of a*
- <s> *seen him </s> (stopped at less than 10 words because a </s> token was encountered)*

Unigram sentence generated from neg.txt:

- *as brown right you of the just drama just lacks*
- *detached -- have exploitive burger disaster of of anyone mr*
- *hugh surrounding director be haplessness as still over ve leniency*

Bigram sentence generated from neg.txt:

- <s> *the most visually flashy camera movements than us only way*
- <s> *the rock s not only fun of disguise falls short*
- <s> *starts and in the genre </s> (stopped at less than 10 words because a </s> token was encountered)*

The unigram sentences appear to be a random selection of words from the corpus, as expected. The bigram sentence more closely resembles a well-formed English sentence. This is because we are able to draw on context when choosing the next word, and so from word to word, the output is closer to the well-formed sentences in the corpus. [addition from original Part 1 which was missing this] When analyzing the positive vs negative generated sentence, there is nothing to indicate a positive or negative sentiment. There are phrases like “the most visually flashy” in the negative generated bigram text that actually sound more positive than negative. Overall, there is no particular trends we can conclude from this, other than that a unigram and bigram model for both positive and negative corpora are quite terrible for generating text on their own, as they can’t capture the intricacies of grammar.

Experimentation with seeding:

We also used seeding to see if better sentences could be generated. We provide a function that will take in the seeded words, and will then append n extra words, using the seeded words as context for the bigram model we have. We have not done this for unigram, as the extra context is not going to be used. Even for bigram, only the last word will be useful, as it is what we use to randomly choose a value.

We notice that with seeding, the sentence structure is still around the same as what the original bigram models provide in terms of English sentences that make sense. This is likely because the word that is used for context to feed into the bigram model is not sufficient to really change the probability of what is returned. If we were to adopt a larger window of context, we would likely see a more dramatic effect between seeding and no seeding on sentence generation.

Bigram generated sentence with seed [simply radiates star-power potential]: *simply radiates star-power potential success </s>*

Bigram generated sentence with seed [simply radiates star-power potential]: *simply radiates star-power potential for something one of the charisma make the dragons !*

Part 4) Smoothing and unknown words

The goal was to implement smoothing and a method for handling unknown words.

For smoothing, we decided to implement +k smoothing. We debated trying +1 smoothing, but because +1 smoothing is generally inadequate, we jumped straight to +k. To find the k values, we implemented a two stage process, where we would first increase k values by 10 (from $k=1 \cdot 10^{-9}$ up to $k=1$), then build a new range around the best k value to narrow it down further.

To evaluate the effectiveness of a k value, we built an updated probability model for them with the training data by summing the count numerator by k and denominator by $k \cdot N$, where N is the number of tokens. We then evaluated the effectiveness of this model on the validation data set, for both positive and negative data sets. We evaluated the effectiveness by adding the negative log probabilities for the validation corpus, and then choosing the k that minimizes the perplexity of the dataset.

The results from this process will be discussed in Part 5), as that is where we calculated the perplexity.

For unknown words, we decided to assign the count for all words that occurred once and only once to the unknown counter, <unk>. This was a standard method discussed in class that we felt would perform sufficiently well. We then calculated the probability distribution model with these updated counts. For unigrams, this was sufficient to include unknown words in the probability mass. We also kept track of which words we'd converted into unknown words. Then for bigrams, if either of the two tokens in the bigrams were in the unknown words set, we replace it (for the purpose of the count, not in the actual data set), and increment the resulting (<unk>, token), (token, <unk>), (<unk>, <unk>) entry accordingly. After this, we then convert the counts into a suitable probability distribution, with +k smoothing as described above.

For unseen bigrams, we had also added a 0 count probability for every unseen combination of word types. This was incredibly inefficient, and unnecessary. Any bigram in the validation/test data set that had both its tokens in the training corpus, but not the combination together, could be assumed to be an unseen bigram, for which would assign a default probability, non-zero because of our +k smoothing. We did not apply unseen ngrams to the unigram, as we couldn't see how we would ever have unseen unigrams. Any unigram that appears in the training corpora will either have a count of 1, and be treated as unknown, or more than 1, and have a count assigned. This is different for bigrams, because we could have the possibility where the individual tokens all appeared in the training corpus, just not in the specific combination.

Part 5) Perplexity

Here, the goal was to implement a mechanism for evaluating the perplexity of our models. We used the formula provided within the project description, with no major modification. Note that the formula described the sums of the conditional probabilities of a token given the rest of the words in that n-gram. This is essentially what we calculated as our language model in Part 3, and enhanced with smoothing in Part 4. With Part 5, we were now able to evaluate the effectiveness of our probability models, smoothed with different values k. The results are displayed below:

Test case	Perplexity	K value
Positive Unigram	1.78	$1 * 10^{-9}$
Positive Bigram	2.01	0.011
Negative Unigram	1.76	$1 * 10^{-9}$
Negative Bigram	1.97	0.011

To test perplexity, we ran through a logarithmic scale from $1 * 10^{-9}$ to 1, increasing by $*10$ every time. Then, with the smallest perplexity value k, we went through a smaller scale centered around k, to try and finetune the optimal k value.

As we see for unigrams, the k value is minimal. This likely indicates that we are overfitting on the data for unigrams, and that our eventual model will not handle unseen tokens well (tokens it has seen before will have a disproportionate effect on the end result). This is likely because unigrams do not have enough context, and so we are shifting over too much probability mass for no reason. While we left it as is to see what impact it would have, we would normally add some nominal k value anyways to avoid overfitting our model.

For bigrams, we observe different results. The optimal k value, when the perplexity of the model is minimized, is at 0.011. This is a good result, as it illustrates the concave curve we expected to see in our dataset. I separated both the k for positive and negative models, even though they ended up being the same. This shows that we shifted over neither too much nor too little probability mass. While we may still be overfitting on our models, it's less likely. I then used this computed k value for Part 6, when I used language models to classify sentences based on their sentiment.

A further enhancement we could have made on this would have been to stagger how we added/removed probability mass. That is, instead of just adding +k across the board, we could have added more probability mass to tokens that showed up rarely, and removed it from tokens that appeared more frequently in the training set.

Part 6) Sentiment Classification

Here, the goal was to perform sentiment classification. When we saw this, our immediate thought process was that in previous parts, we'd already built separate probability models for positive and negative corpora, and that we could leverage those here. We decided to adopt a similar approach to how we'd handled perplexity. We first split up the (validation and test) corpora into sentences, and preprocessed as necessary. We'd decided to test it with unigrams first, so we simply fed each token into our probability model, then summed up the negative log of the ensuing product. We then compared the outputs from positive and negative probability models, and classified the sentence based on which had the higher resulting probability. We then repeated the same process with bigrams instead of unigrams.

We then submitted two runs to Kaggle. The first was a unigram model with +1 smoothing, the simplest test case, just to make sure that our process was working. The second was the bigram model with +k smoothing we had developed. To our surprise, it had a worse performance than unigram model with +1 smoothing. We will discuss why after reviewing the outcome from our validation data set.

Results from Kaggle:

Run Type	
Unigram model with +1 smoothing	0.58043
Bigram model with +k smoothing	0.55924

We also tested against our validation data sets. The results are below.

When we treat every 1 count word as unk:

- Expected all sum_pos_uni, but got 299 sum_pos_uni and 118 sum_neg_uni
- Expected all sum_pos_bi, but got 277 sum_pos_bi and 140 sum_neg_bi
- Expected all sum_neg_uni, but got 206 sum_pos_uni and 202 sum_neg_uni
- Expected all sum_neg_bi, but got 207 sum_pos_bi and 201 sum_neg_bi

When we only treat every second 1 count word as unk:

- Expected all sum_pos_uni, but got 285 sum_pos_uni and 132 sum_neg_uni
- Expected all sum_pos_bi, but got 267 sum_pos_bi and 150 sum_neg_bi
- Expected all sum_neg_uni, but got 240 sum_pos_uni and 168 sum_neg_uni
- Expected all sum_neg_bi, but got 217 sum_pos_bi and 191 sum_neg_bi

When we only treat every third 1 count word as unk:

- Expected all sum_pos_uni, but got 269 sum_pos_uni and 148 sum_neg_uni
- Expected all sum_pos_bi, but got 258 sum_pos_bi and 159 sum_neg_bi
- Expected all sum_neg_uni, but got 214 sum_pos_uni and 194 sum_neg_uni
- Expected all sum_neg_bi, but got 203 sum_pos_bi and 205 sum_neg_bi

When we treat every fourth 1 count word as unk:

- Expected all sum_pos_uni, but got 258 sum_pos_uni and 159 sum_neg_uni
- Expected all sum_pos_bi, but got 245 sum_pos_bi and 172 sum_neg_bi
- Expected all sum_neg_uni, but got 227 sum_pos_uni and 181 sum_neg_uni
- Expected all sum_neg_bi, but got 203 sum_pos_bi and 205 sum_neg_bi

Looking at our results, we observed some very strange trends. For positive test inputs, our unigram model performed better than our bigram model. For the negative model, it performed worse than the bigram, and even than a coin toss. For the Kaggle data, we couldn't distinguish positive and negative input, but we could see that the unigram model with +1 smoothing, our initial test point, was marginally better than bigram model with +k smoothing. At this point, we evaluated how we were assigning counts to unknown, and decided to reduce the number of unknown words by only assigning every second, third or fourth 1 count token to be unknown. This did not help the performance of the system measurably, so we kept using every unknown word

When looking at why this was the case, we came up with a few theories. First, our data set is very small. This could cause outliers in the training data to have an adverse effect on our model. In addition, we did not remove filler words as part of our preprocessing. At the time, we thought it would not be necessary, but now, looking at the results, we suspect it had an adverse effect.

When we referenced literature to find what a "good" sentimentality result would have been, Roebuck (2012) showed that a program with 70% accuracy is actually close to how a human performs. This indicates that our positive unigram probability model, where we treat every 1 count token as *<unk>* is actually performing very well, and our negative probability model, for both unigram and bigram and across all test cases is not. We could not identify a specific reason, other than the dataset itself, for why the performance was so poor. A future improvement could be to also expand from bigrams to trigrams or quadgrams, in the hope that this extra context would capture more accurately the sentiment of the statement.

Part 7) No work was done for this; we chose Part 8) instead

Part 8.1) A machine learning variant of sentiment classification with word embedding

We chose part 8.1 among the alternatives available because we were familiar with ML strategies and this problem appealed to us more than the others – we wanted to see how the ML algorithms we had learnt could be applied to NLP and text classification, and also experiment with word embeddings.

The high level approach that we followed for this task was to:

- Come up with a good feature representation of the movie reviews in the training data
- Train a Machine Learning model that gives good accuracy on the 'dev' dataset
- Use the 'dev' dataset again to perform cross-validation to tune the hyper parameters of our model
- Finally make predictions using the tuned model on the test data.

We tried a couple of ML algorithms that we were familiar with and that we expected to do well for the binary classification task we were given. We decided to start with linear classifiers, as they are not only simple, but also well suited for the initial problem, and we were expecting an acceptable accuracy.

We therefore first tried with Naïve Bayes and simple SVMs. The results are below:

ML Algorithm	Accuracy
Naïve Bayes before word embedding	51.42%
Naïve Bayes with word embeddings	69.09%
LinearSVC with word embeddings	72.72%
RBF	81.28%

Naïve Bayes without word embeddings was little better than a coin toss. For coming up with a good feature representation of the movie reviews, we first used the Word2Vec functionality in Python from the "gensim" library to create the word embeddings ourselves from the given corpus. This was giving us very bad accuracy, as SVM was not able to learn a good enough decision boundary. After looking at the discussions on Piazza, we realized our mistake on the way we were generating features for the model, and decided to experiment with some of the pre-trained word embedding suggested in the problem document. We used the pre-trained word vectors from Google News dataset which has 300-dimensional vectors for 3 million words and phrases. This significantly improved the accuracy of Naïve Bayes and LinearSVC which we also implemented, and we did not experiment with the other pre-trained embeddings like Glove.

We felt that we could do better with non-linear models by using kernels. We knew that Radial Basis Function Kernels (RBF Kernel) are a universal approximator and work well in most scenarios. After playing around a bit with the "C" value which determines the amount of slack for allowing points within the margins of the SVM, we were able to train a classifier that gives us 81.28%. We determined the ideal value for "C" using cross-validation on the

“dev” dataset on which we were getting 72.48 % accuracy(max). We used K-fold Cross Validation with 10 folds (K=10). While more splits could have been possible, it would have led to more iterations and in-turn significantly higher validation time.

Distribution of work:

Vinayaka kicked off the project with the pre-processing, followed by which Nithish computed the unsmoothed unigram and bigram probabilities. Once this was completed, Akash was able to independently work on choosing the optimal ‘K’ for smoothing. Part 5, 6 and 8.1 were done completely parallelly.

Akash Peri (ap659): Owned part 5 and part 6 and collaborated on smoothing

Vinayaka (vb247): Owned part 2, part 4 and helped in tuning parameters for smoothing

Nithish (nr293): Owned part 3 and part 8.1

References:

Roebuck, K. (2012-10-24). Sentiment Analysis: High-impact Strategies - What You Need to Know: Definitions, Adoptions, Impact, Benefits, Maturity, Vendors